

AI-Assisted Video Pre-Screening Platform

Turning first-round hiring from a scheduling-heavy, subjective bottleneck into a fast, consistent, and fully explainable shortlisting engine — with a human always making the final call.

INDUSTRY	FOCUS	APPROACH	PLATFORM
Enterprise Software	Talent Acquisition · Hiring	Human-in-the-Loop AI	Cloud-Native · Multi-Provider AI

Confidentiality note. This is a sanitized portfolio overview. The client identity, product name, proprietary rubrics, and internal source are withheld under NDA. Everything here describes capabilities and engineering approach at a level safe for public sharing. Outcome figures are representative and anonymized.

© 2026 Rohail K. Malhi — All rights reserved. Proprietary work; not open-source. The frameworks, architecture, metrics, and artifacts shown here may not be copied, adapted, or repurposed. Shared solely to evaluate the author's experience.

Screening doesn't scale — and it isn't fair when it's rushed

For a single open role, a recruiter may screen dozens or hundreds of applicants through 20–30 minute live calls, most ending in a "no." The cost compounds well beyond the hours spent:

- **Scheduling friction** — coordinating live calls across calendars and time zones stalls the funnel for days per candidate.
- **Inconsistent evaluation** — different interviewers, days, and questions mean early and late candidates are judged on different scales.
- **Subjectivity & bias risk** — unstructured "gut-feel" screening is hard to defend and raises fairness and compliance concerns.
- **No audit trail** — when a decision is challenged, there's rarely a record of *why* a candidate advanced.
- **A hard throughput ceiling** — capacity is capped by how many live calls a recruiter can physically run.

The client needed to screen **more candidates, faster, and more fairly** — without adding headcount, and without replacing human judgment with a black box.

Asynchronous video interviews with explainable, human-in-the-loop AI

Candidates answer a structured set of questions on their own time via short, single-take video responses. The platform transcribes and analyses each answer across multiple dimensions, then produces an **explainable recommendation** — a score, a confidence level, and a ranking — that a human reviewer uses to decide who advances. The guiding principle throughout: **AI assists, a human decides.**

The candidate experience

- **Frictionless, no-account access** via a unique, secure, expiring invite link.
- **Informed consent first** — an explicit, timestamped step before any recording.
- **Structured, timed, single-take answers** that auto-submit — authentic and comparable, not over-rehearsed.
- **Desktop and mobile**, with a guided device check and a low-pressure practice question.

The evaluation engine — three lenses, each in its proper role

● **Content & substance** · scored

Rubric-based scoring with criterion-by-criterion reasoning, STAR-structure checks, semantic match to an ideal answer, and surfacing of contradictions, resume mismatches, likely AI-generated text, and red-flag statements.

● **Communication & delivery** · advisory

Objective, transparent measures of pace, pauses, fluency, and language — shown as context for reviewers, never folded into the score.

● **Integrity** · flags

Presence and session-consistency checks and detection of off-camera prompting — surfaced as flags for a person to weigh, never as an automated verdict.

● **Recruiter-ready output**


An overall fit score, a confidence value that tells reviewers where to focus, a ranked shortlist, per-question highlights with justifications, and every flag — all traceable to transcript and video.

The reviewer experience

- A **ranked candidate dashboard** per role, sorted by fit and confidence.
- A **per-candidate deep-dive**: playback, transcript, delivery metrics, integrity flags, and rubric scores *with their justifications*.
- **Human override, always** — adjust any score or ranking and record the decision with notes; every override is audit-logged.

Decoupled, asynchronous, provider-agnostic

Engineered so heavy AI work never blocks candidates, and so the AI layer can flex across providers without touching the product.




End-to-end flow — capture never depends on AI availability; the queue absorbs all pipeline backpressure.

LAYER	STACK
Frontend	React · TypeScript · Vite (single-page app; in-browser media capture)
Backend	Python · FastAPI (JSON API + background worker)
Data	PostgreSQL · durable job queue with retry & idempotency
Storage	S3-compatible object storage · presigned direct upload
AI orchestration	Provider-agnostic gateway — rate limiting, token budgeting, automatic failover, quota visibility
AI capabilities	Speech-to-text with diarization & word-level metadata · LLM rubric scoring with structured output · semantic embeddings · vision-based presence/integrity signals
Delivery	Cloud-hosted services · environment-isolated credentials · privacy-aware data handling

SOLUTION · EVALUATION MODEL

How a candidate is evaluated

Three independent lenses, combined transparently — never a single opaque score.



RESULT

From a weeks-long backlog to same-day shortlists

The platform reframes first-round screening from a serial, calendar-bound activity into a parallel, always-on one — while making every decision more consistent and more defensible.

~70%

LESS SCREENING TIME

Async answers replace live calls

3–4×

MORE CANDIDATES /
RECRUITER

No added headcount

**2 wks → 3
days**

TIME-TO-SHORTLIST

Typical requisition

100%

RUBRIC CONSISTENCY

Every candidate, same scale

0

ROUND-1 SCHEDULING

Eliminated entirely

Full

AUDIT TRAIL

Every recommendation
explainable

"We went from a screening backlog measured in weeks to same-day shortlists — and for the first time we can actually show why each candidate ranked where they did."

— Head of Talent Acquisition, enterprise software client (identity withheld under NDA)

RESPONSIBLE AI & ENGINEERING

Built to be trusted, not just clever

- **Human-in-the-loop by construction.** The system never auto-rejects or auto-hires — it recommends; people decide, and overrides are first-class and audited.
- **Signal discipline.** Delivery metrics are advisory and never drive the score, reducing the risk of penalizing accent, speech pattern, or delivery style.
- **Explainable, not oracular.** Every score carries its justification and traces back to the transcript and recording.
- **Privacy by default.** Explicit, timestamped consent precedes any capture; data handling is built around minimization, controlled retention, and terms suitable for sensitive personal data.
- **Fault-tolerant AI orchestration.** Rate-limit-aware, multi-provider failover and idempotent, retryable stages keep the pipeline running through provider quotas and outages.
- **Cost-aware architecture.** An asynchronous, provider-agnostic design keeps infrastructure lean and lets the AI layer be tuned for cost and quality without product changes.

At a glance — an asynchronous, AI-assisted video pre-screening platform that lets talent teams screen far more candidates, far faster, and far more consistently: assessing each answer for content, delivery, and integrity; producing an explainable, ranked, confidence-scored shortlist; and keeping a human firmly in control of every hiring decision, with a full audit trail behind it.